

Introduction to Machine Learning and Deep Learning

Plan

- 25/04 Introduction
- 02/05 Calculus and Linear Algebra
Linear models and logistic regression
- 09/05 SVM; k-fold cross-validation and boosting
- 16/05 CNNs; Backprop; Representation Learning; Regularisation; SGD
- 23/05 Image classification using Deep Learning models;
Keras, Tensorflow and TF-tensorboard

Feedback is welcome

Expect changes and corrections

Please be aware it is not the same room all the times



theano **PYTORCH**

(and many, many more)



IP[y]:
IPython



Python

- Numpy, scipy
- Scikit-learn
- Scikit-image, opencv, pillow
- Pandas
- Tensorflow, theano, keras, pytorch, mxnet, cntk
- Matplotlib, bokeh

IDE: Jupyter notebook, Spyder, PyCharm

DL framework requirements

- Need to work with tensors
- Neural Network primitives
- Symbolic derivatives to do efficient backprop
- Offloading to GPUs
- Support for parallelism (multi-device)

Frameworks

- **Caffe, Caffe2** - Berkeley, FB; model Zoo
- **Theano** - U. Montreal; Python, symbolic derivatives
- **TensorFlow** - Google; Python, symbolic derivatives
- **Pytorch** - FB; Python, Dynamic Neural Networks
- **Apache Mxnet** – Amazon (+ Intel, Baidu and others); Python
- **Microsoft Cognitive Toolkit (CNTK)** - Microsoft
- **Keras** - François Chollet, Python, runs atop of TF, Theano
- **Lasagne** - Python, runs on Theano

Setting yourself up for hands-on experiments⁷

- When using laptop we suggest to install Anacoda or miniconda
<https://conda.io/miniconda.html>
<https://www.anaconda.com/download/>
- You can also use Davros and setup your environment there
- Further details on setting up conda environment
<https://shadow.icr.ac.uk/Wiki/conda>
- One easy option is MS Notebook service on Azure which is free but you only get two vCPUs
- Remember, serious training requires one or more GPUs

Running ipython notebooks on MS Azure:

1) Go to <https://notebooks.azure.com/>

Microsoft Azure Notebooks Preview

Sign In

Libraries What's New Help



Now Featuring: Python Data Science Handbook by Jake VanderPlas

Interactive coding in your browser

Free, in the cloud,
powered by Jupyter


Get Started

Plot Iris data using matplotlib/seaborn

```
In [1]: %matplotlib inline
import seaborn as sns
sns.set()

In [2]: iris = sns.load_dataset("iris")
g = sns.lmplot(x="sepal_length", y="sepal_width", hue="species",
              truncate=True, size=5, data=iris)
g.set_axis_labels("Sepal length (mm)", "Sepal width (mm)")


Out[2]: <seaborn.axisgrid.FacetGrid at 0x7f0e0ff0fdd0>
```



Ask me anything

2) Sign in with any MS account

9



We have a new sign-in experience!
[Try it now](#)



Azure Notebooks

Work or school, or personal Microsoft account

☐ Keep me signed in

Sign in

[Can't access your account?](#)

[Other sign-in options](#)

[Create a new Microsoft account](#)

3) Add new library "Introduction to Python": + New Library

The screenshot shows the 'Create New Library' dialog box in the Ink1 Libraries application. The dialog has a title bar with a question mark and a close button. It features two tabs: 'New' (selected) and 'From GitHub'. The 'Library Name' field contains 'Introduction to Python'. The 'Library ID' field is split into two parts: 'ink1/libraries/' and 'IntroductionToPython'. There are two checked checkboxes: 'Public library' and 'Create a README.md'. At the bottom right, there are 'Create' and 'Cancel' buttons. The background shows the 'Libraries' page with a sidebar containing a search bar and a list of libraries: 'basic python', 'cntk', 'pytorch tutorial', and 'Sample notebooks'. The status 'Showing 4 libraries' is at the bottom left, and a pagination control with '< 1 >' is at the bottom right.

Libraries What's New Help

ink1 > Libraries

+ New Library

Search

NAME ▼

- basic python
- cntk
- pytorch tutorial
- Sample notebooks

Showing 4 libraries

< 1 >

Stopped 0 Apr 29, 2017 May 2, 2017

Contact us FAQ Privacy and cookies Terms of use Trademarks © 2017 Microsoft

4) Start Python notebook: + New

The screenshot shows the Microsoft Azure Notebooks web interface. The browser address bar displays the URL `https://notebooks.azure.com/ink1/libraries/IntroductionToPython`. The page header includes the Microsoft Azure Notebooks logo and the user name 'ink1'. A navigation bar contains links for 'Libraries', 'What's New', 'Status', and 'Help'. The main content area features a large blue banner with the text 'Introduction to Python' and an icon of three books. Below the banner, a breadcrumb trail shows 'ink1 > Libraries > IntroductionToPython'. A toolbar at the top of the workspace contains various actions: Run, New, Settings, Share, Clone, 0 Clones, Star (0), Terminal, Shutdown, Preview, Edit File, Download, and Delete. Below the toolbar is a search bar with the text 'Search' and a button 'Show hidden items'. A table lists the files in the workspace, with columns for 'FILE NAME', 'FILE TYPE', and 'MODIFIED'. The table contains one entry: 'README.md' (Markdown) modified on 'Apr 13, 2018'. Below the table, it says 'Showing 1 file'. At the bottom of the page, there is a footer with links for 'Contact us', 'FAQ', 'Privacy and cookies', 'Terms of use', 'Trademarks', and a copyright notice '© 2017 Microsoft'.

Microsoft Azure Notebooks Preview ink1

Libraries What's New Status Help

Introduction to Python

ink1 > Libraries > IntroductionToPython

Run + New Settings Share Clone 0 Clones ☆ Star (0) Terminal Shutdown Preview Edit File Download Delete

Search Show hidden items

FILE NAME ▼	FILE TYPE	MODIFIED
README.md	Markdown	Apr 13, 2018

Showing 1 file

< 1 >

Contact us FAQ Privacy and cookies Terms of use Trademarks © 2017 Microsoft

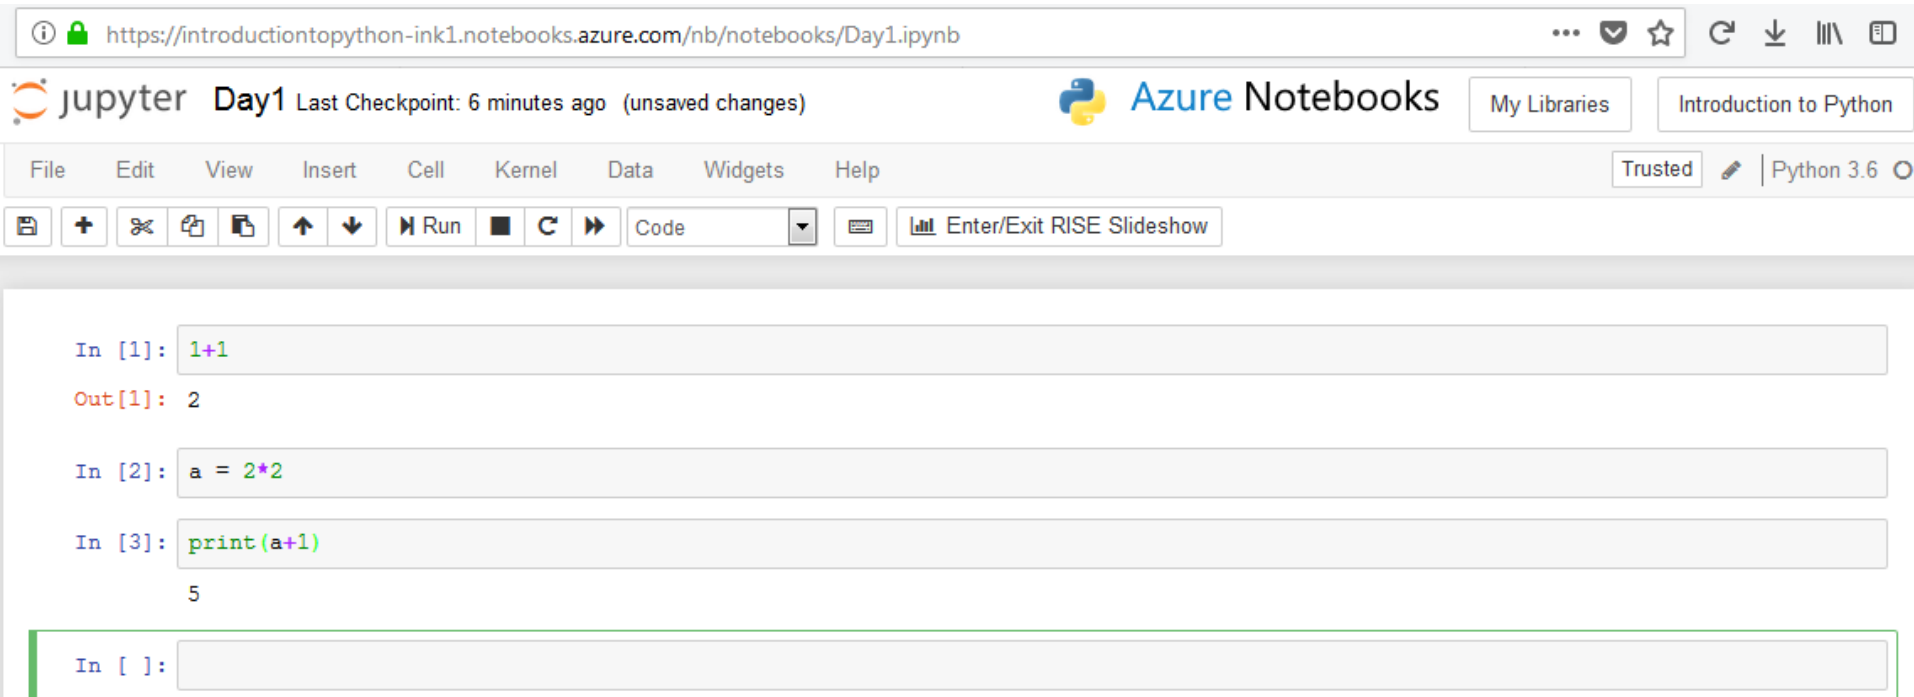
5) Start Python notebook: select Python version

The screenshot shows the Microsoft Azure Notebooks web interface. A modal dialog titled "Add Items to Library" is open in the center. The dialog has three tabs: "New", "From computer", and "From URL". The "New" tab is selected. Inside the dialog, there is a text input field for "Item Name" containing the text "Day1". Below it is a dropdown menu for "Item type" with the following options: "Blank File", "Python 2.7 Notebook", "Python 3.5 Notebook", "Python 3.6 Notebook" (which is highlighted in blue), "R Notebook", "F# Notebook", and "Folder". At the bottom right of the dialog are two buttons: "New" and "Cancel".

The background interface shows the URL <https://notebooks.azure.com/ink1/libraries/IntroductionToPython>. The page title is "Microsoft Azure Notebooks". The breadcrumb navigation is "ink1 > Libraries > IntroductionToPython". The file list shows one file: "README.md" with a file type of "Markdown" and a modification date of "Apr 13, 2018". The footer contains links for "Contact us", "FAQ", "Privacy and cookies", "Terms of use", "Trademarks", and a copyright notice "© 2017 Microsoft".

6) start using your notebook

13



The screenshot displays the Azure Jupyter Notebook interface. The browser address bar shows the URL: <https://introductiontopython-ink1.notebooks.azure.com/nb/notebooks/Day1.ipynb>. The notebook title is "Day1" with a subtitle "Last Checkpoint: 6 minutes ago (unsaved changes)". The Azure Notebooks logo is visible. The top menu bar includes File, Edit, View, Insert, Cell, Kernel, Data, Widgets, and Help. The right side of the menu bar shows "Trusted" and "Python 3.6". The toolbar contains icons for saving, adding, deleting, and running cells, as well as a dropdown menu for cell types (Code, Text, Raw) and a button for "Enter/Exit RISE Slideshow". The main area shows three code cells. The first cell contains `1+1` and has an output of `2`. The second cell contains `a = 2*2`. The third cell contains `print(a+1)` and has an output of `5`. A fourth, empty code cell is at the bottom, highlighted with a green border.

https://introductiontopython-ink1.notebooks.azure.com/nb/notebooks/Day1.ipynb

jupyter Day1 Last Checkpoint: 6 minutes ago (unsaved changes) Azure Notebooks My Libraries Introduction to Python

File Edit View Insert Cell Kernel Data Widgets Help Trusted Python 3.6

Run Code Enter/Exit RISE Slideshow

```
In [1]: 1+1
Out[1]: 2

In [2]: a = 2*2

In [3]: print(a+1)
5

In [ ]:
```

7) adding a file to your library

14

https://notebooks.azure.com/ink1/libraries/IntroductionToPython

Microsoft Azure Notebooks

ink1 > Libraries > IntroductionToPython

Libraries What's New Status

Introduction

ink1 > Libraries > IntroductionToPython

Run + New Settings

Search Show hidden items

FILE NAME

Day1.ipynb

README.md

test.py

Showing 3 search results (1 hidden)

Running

Edit File Download Delete

FILE NAME	FILE TYPE	MODIFIED
Day1.ipynb	Notebook	Apr 15, 2018
README.md	Markdown	Apr 14, 2018
test.py	Python	Apr 15, 2018

< 1 >

Add Items to Library

New From computer From URL

Item Name

file.py

Item type

- Blank File
- Blank File
- Python 2.7 Notebook
- Python 3.5 Notebook
- Python 3.6 Notebook
- R Notebook
- F# Notebook
- Folder

New Cancel



8) editing a file - right click - Edit File

https://notebooks.azure.com/ink1/libraries/IntroductionToPython

Microsoft Azure Notebooks Preview

ink1

Libraries

What's New

Status

Help

Introduction to Python

ink1 > Libraries > IntroductionToPython

> Running

Run

+ New

Settings

Share

Clone

0 Clones

Star (0)

Terminal

Shutdown

Preview

Edit File

Download

Delete

Search

Show hidden items

FILE NAME	FILE TYPE	MODIFIED
Day1.ipynb	Notebook	Apr 15, 2018
file	Python	Apr 15, 2018
R	Markdown	Apr 14, 2018
te	Python	Apr 15, 2018

Show

hidden

Run (r)

Preview (p)

Edit File (i)

Download (d)

Rename (a)

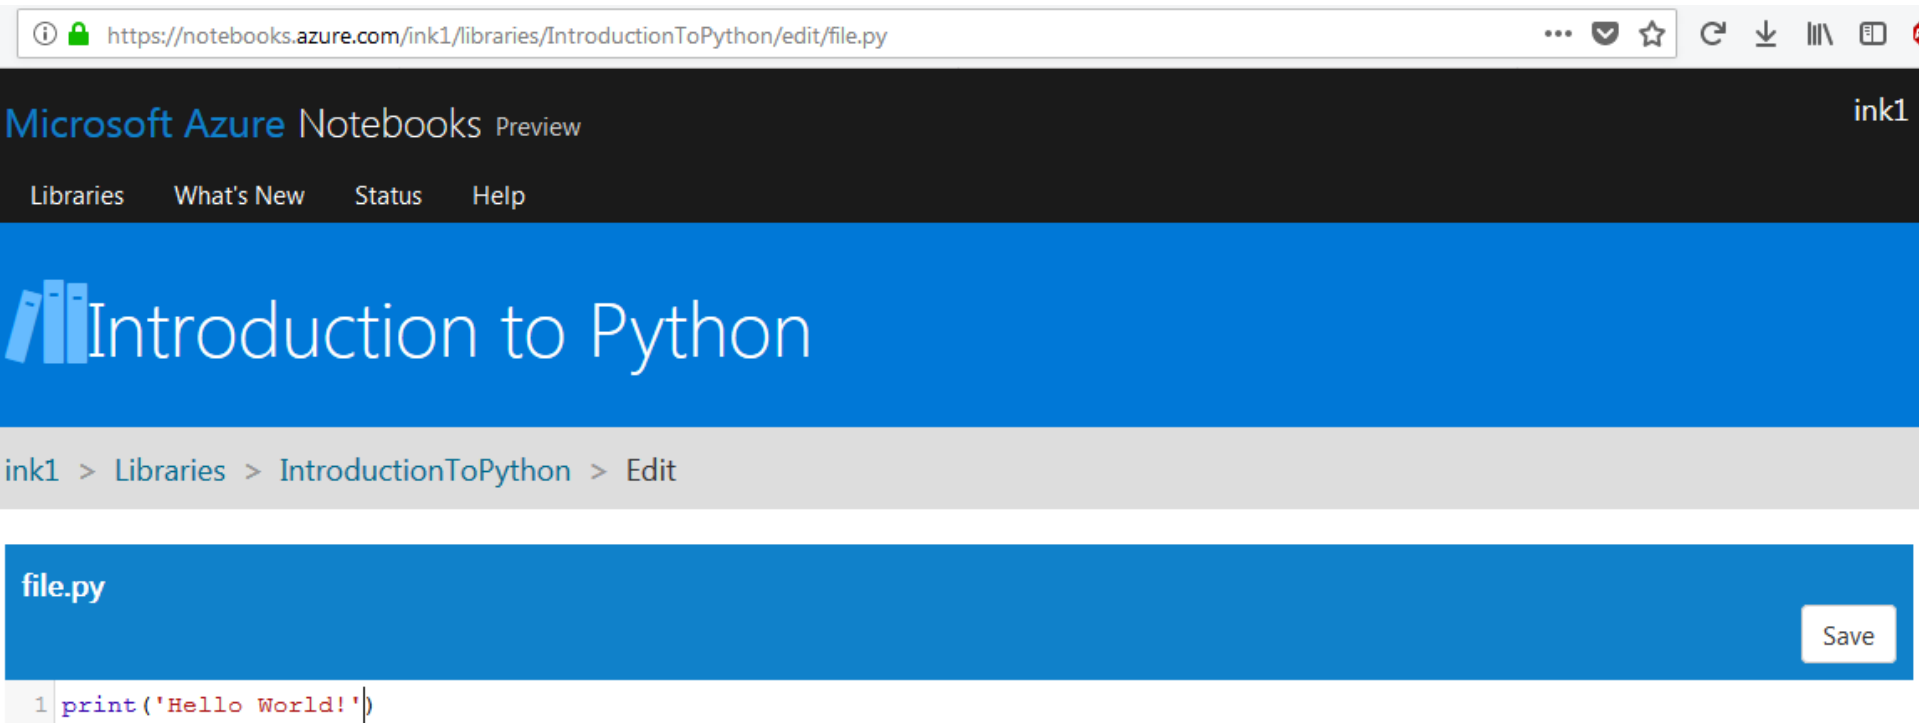
Delete (x)

<

1

>

9) add one line
`print('Hello World!')`
and save the file



The screenshot shows the Microsoft Azure Notebooks interface. At the top, there's a browser address bar with the URL `https://notebooks.azure.com/ink1/libraries/IntroductionToPython/edit/file.py`. Below the address bar, the header displays "Microsoft Azure Notebooks Preview" and the notebook name "ink1". A navigation bar includes links for "Libraries", "What's New", "Status", and "Help". The main content area has a blue header with the text "Introduction to Python" and a breadcrumb trail: "ink1 > Libraries > IntroductionToPython > Edit". Below this, a blue bar indicates the file name "file.py" and contains a "Save" button. The code editor shows a single line of Python code: `1 print('Hello World!')`.

10) run the file

17

https://introductiontopython-ink1.notebooks.azure.com/nb/notebooks/Day1.ipynb

jupyter Day1 Last Checkpoint: 19 minutes ago (unsaved changes) Azure Notebooks My Libraries Introduction to Python

File Edit View Insert Cell Kernel Data Widgets Help Trusted Python 3.6

Run Enter/Exit RISE Slideshow

```
In [1]: 1+1
Out[1]: 2

In [2]: a = 2*2

In [3]: print(a+1)
5

In [5]: run file.py
Hello World!

In [ ]:
```

Conda package management

```
# open terminal through Azure notebook
# this gives you access to the VM itself

nbuser@nbserver:~$ . anaconda3_501/etc/profile.d/conda.sh
nbuser@nbserver:~$ conda list

...

nbuser@nbserver:~$ which python
/usr/bin/python

nbuser@nbserver:~$ conda activate base
(base) nbuser@nbserver:~$ which python
/home/nbuser/anaconda3_501/bin/python
(base) nbuser@nbserver:~$ conda install -c damianavila82 rise
```

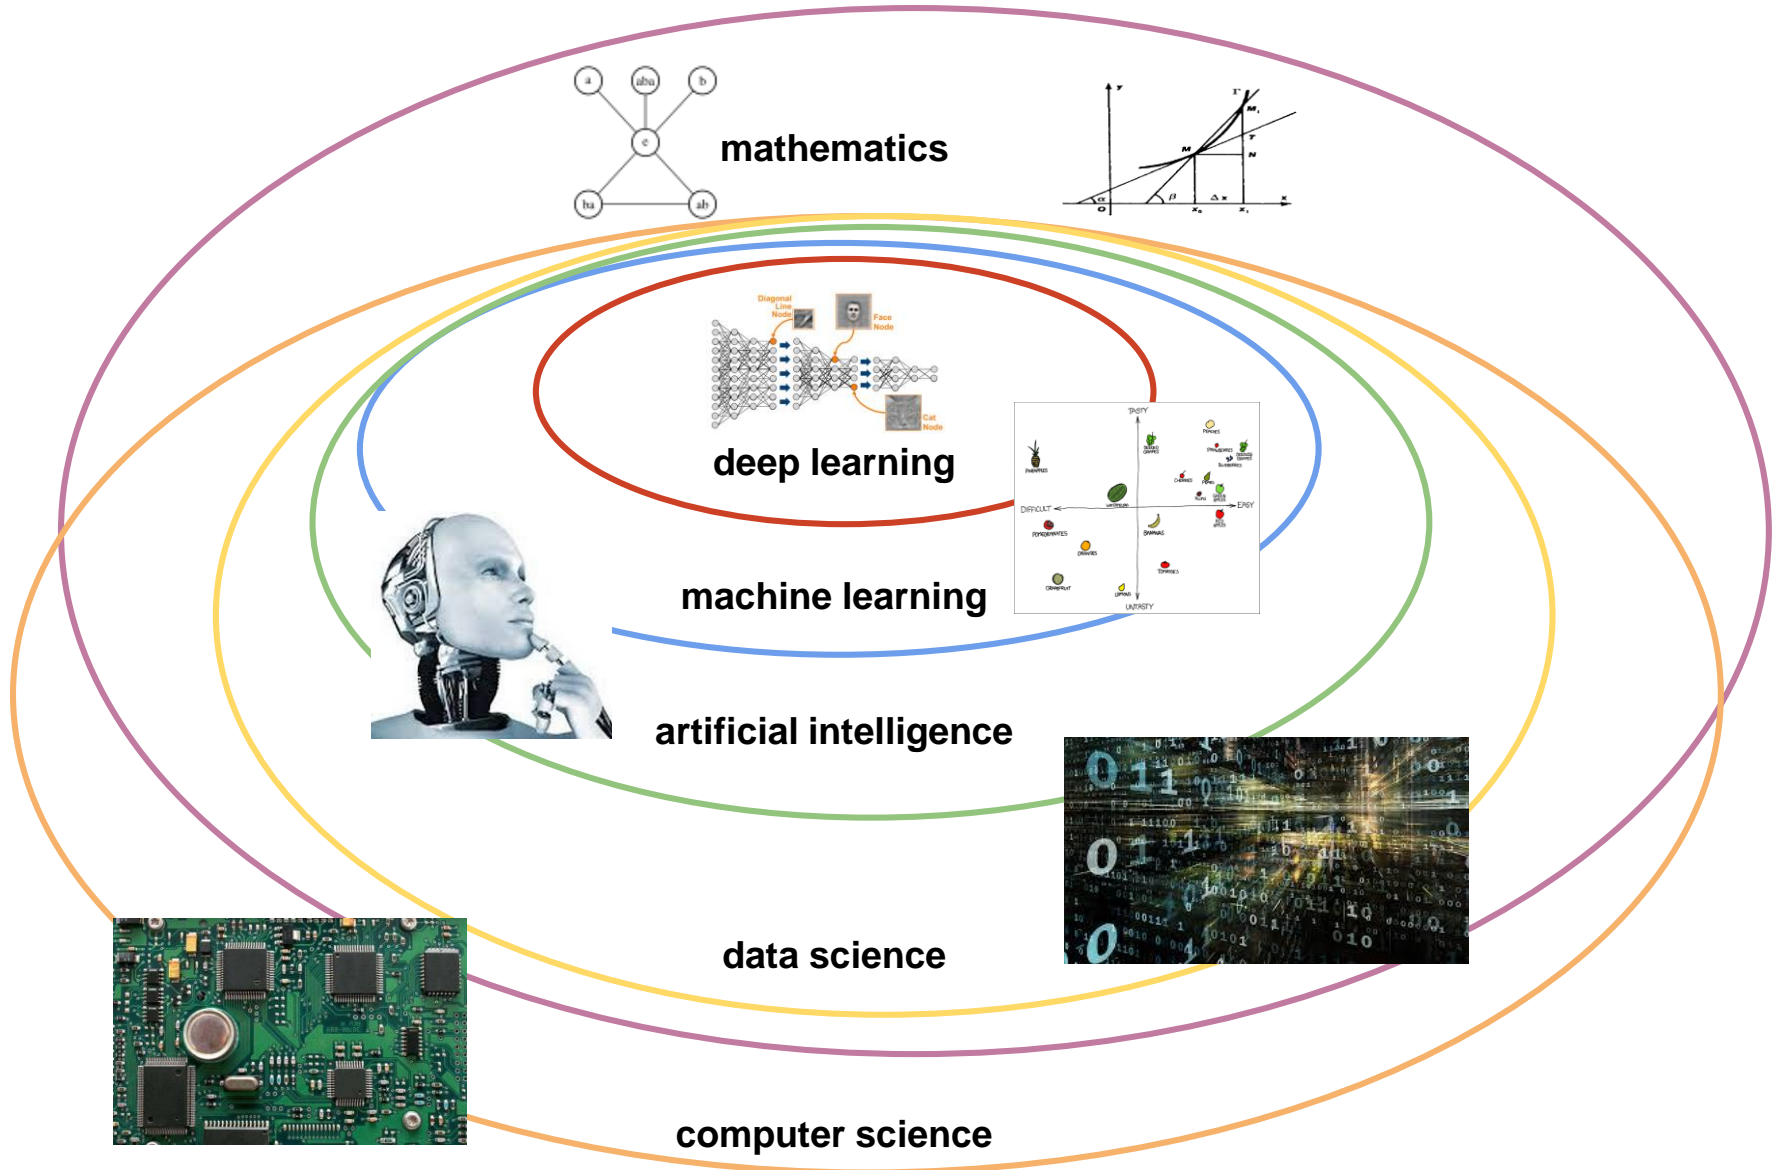
Conda package management

```
# originally an alternative Python & env manager
# available as Conda or Anaconda (conda + ~200 packages)
# https://www.continuum.io/downloads

~> conda create --name myenv keras
~> conda activate myenv
(myenv) ~> conda list --export > package-list.txt
(myenv) ~> conda deactivate
~> conda info -e          # list environments
```

Deep Learning Universe

20



Definitions



Artificial Intelligence (AI) born in 1950s, the effort to automate intellectual tasks normally performed by humans.

Symbolic AI, expert systems of 1980s

Explicit rules, logical problems such as playing chess

Machine Learning (ML) is learning from data as an alternative to manual coding for specific response, features etc.

Deep Learning (DL) is a branch of ML employing multi-layered Neural Networks (NNs). The word “deep” is used to differentiate from “shallow” NNs employing a few layers only.



Machine Learning

- Model is trained rather than programmed
- ML finds statistical structure in data that allows the system to come up with rules for automating the task
- ML is essentially a form of applied statistics with increased emphasis on the use of computers to statistically estimate complicated functions and a decreased emphasis on proving confidence intervals around these functions

Tasks solvable by ML

- Classification $\mathcal{F}: \mathbb{R}^n \rightarrow \{1, \dots, k\}$
 - Classification with missing input
- Regression $\mathcal{F}: \mathbb{R}^n \rightarrow \mathbb{R}$
- Structured output
 - Transcription
 - Translation
 - Image segmentation
- Anomaly detection
- Imputation of missing values $\mathcal{F}: \mathbb{R}^n(\text{with missing values}) \rightarrow \mathbb{R}^n$
- Denoising $\mathcal{F}: \mathbb{R}^n \rightarrow \mathbb{R}^n$

Types of ML Algorithms

Supervised Learning

Learn from training data with labels

Unsupervised Learning

Learn from unlabelled data

(categorize unlabelled data based on discovered similarities)

Reinforcement Learning

Learn from interactions and rewards from the world

(train desired behaviour within a specific context)

Supervised learning in general

- Input data $X = \{X_1, \dots, X_N\}, X_i \in \mathbb{R}^n$
- Output data $Y = \{Y_1, \dots, Y_N\}, Y_i \in \mathbb{R}^m$
- Find mapping $\mathcal{F}(X_i; \omega) = Y_i'$
- such that the loss function $\mathcal{G}(\omega) = \sum_i G(Y_i, Y_i'; \omega)$
- is minimal over the parameter space $\omega \in \mathbb{R}^k$

How to select the right $\mathcal{F}(X; \omega)$?

- Model selection (functional space) – restrictions on the class of functions
- Parameter optimisation – find the best point $\omega \in \mathbb{R}^k$
- Criteria for being the best are far from obvious

Machine Learning

$$Y = \{black, red\}, X \in \mathbb{R}^2$$

Linear transformation in 2D plane allows easy selection of red dots by choosing $x'_1 > 0$.


What does the learning do?

28

Given input data and expected output, compute a representation which can be used to predict the output.

An example of a representation can be HSV (Hue Saturation Value) which is another (exact) representation of RGB format.


The power of a representation is related to model capacity and will in general incur some data loss.



What does the learning do?

29

- All ML algorithms consist of automatically finding such series of transformations that turn input data into more useful representation for a given task. These can be linear transformations, projections, non-linear operations and so on but the hypothesis space is fixed.
- If such a transformation cannot be found immediately the process can proceed iteratively using a feedback from minimisation of a loss function which characterises how close (or far) the current output is from the desired output.



Consider 1D regression

Statistics refresh

31

Discrete random variable X takes values from a set, e.g. coin flip $\{0,1\}$.

Probability of an event $P(A)$.

If $X \in \{0, 1\}$ then $P(X = 0) + P(X = 1) = 1$.

Union of events

$$P(A \vee B) = P(A) + P(B)$$

Joint probability

$$P(A, B) = P(A \wedge B) = P(A|B)P(B)$$

Conditional probability (provided that $P(B) > 0$)

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

Bayes rule

$$P(A|B) = \frac{P(A)P(B|A)}{P(B)}$$

where $P(B)$ may be computed as $P(B) = \sum_a P(B|a)P(a)$

Mean value of probability distribution over all possible events

$$\mathbb{E}[X] = \sum_a a P(a)$$

Statistics refresh

32

Continuous random variable X can take values from \mathbb{R} .

Probability density function $p(x)$.

Probability is $P(a < X < b) = \int_a^b p(x) dx$.

Mean, or expected value μ

$$\mathbb{E}[X] = \int x p(x) dx$$

Variance, measure of “spread” of a distribution, σ^2

$$\text{var}[X] = \mathbb{E}[(X - \mu)^2] = \int (x - \mu)^2 p(x) dx = \mathbb{E}[X^2] - \mu^2$$

Standard deviation

$$\text{std}[X] = \sqrt{\text{var}[X]}$$

Normal distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Statistics refresh

Self-information of an event $X = x$

$$I(X = x) = -\log P(X = x)$$

Shannon entropy

$$H(X) = -\sum_{k=1}^K P(X = k) \log P(X = k) = \mathbb{E}_{X \sim P}[I(X)]$$

Cross entropy of two probability distributions P and Q over the same random variable X

$$H(P, Q) = -\mathbb{E}_{X \sim P}[\log Q] = -\sum_{k=1}^K P \log Q$$

Kullback-Leibler (KL) divergence or relative entropy is a measure of dissimilarity of two probability distributions P and Q over the same random variable X

$$\text{KL}(P||Q) = -H(P) + H(P, Q)$$

Note that generally $\text{KL}(P||Q) \neq \text{KL}(Q||P)$

$\text{KL}(P||Q) = 0$ if and only if $P = Q$

Frequentist Statistics

- Randomness is objective
- Probability is the frequency of past events; in this way it's objective and doesn't depend on one's beliefs
- Data X is random but model parameters ω are fixed (but unknown)
- **Bias** of ω' is the difference between the expectation value of ω' and the true value of ω ; the expectation is over the data X seen as samples from a random variable
- **Variance** of ω' is a measure of how the estimate of ω would change if we resample the dataset from the underlying data generating process
- Ideally we want both bias and variance of our estimate be small and $\omega' \rightarrow \omega$ as $N \rightarrow \infty$
- Find a point estimate ω' using maximum likelihood:
$$\omega' = \arg \max_{\omega} P(X|\omega) = \arg \max_{\omega} \log P(X|\omega)$$

Bayesian Statistics

- Randomness is subjective
- Probability is a measure of our belief so that probability is subjective and refers to the future
- Data X is fixed (given) but model parameters ω are random; search for probability distribution which generates (random) parameters

- Compute posterior

$$P(\omega|X) = \frac{\overset{\text{likelihood}}{P(X|\omega)} \overset{\text{prior}}{P(\omega)}}{\underset{\text{evidence}}{P(X)}}$$

- However most interesting models are intractable, hence point estimates
- Maximum A Posteriori (MAP) point estimate

$$\omega_{MAP} = \arg \max_{\omega} P(\omega|X) = \arg \max_{\omega} \log P(X|\omega) + \log P(\omega)$$

Coin tossing example (after Geoffrey Hinton)

- Suppose we only know each tossing event produces a head with some unknown probability P and a tail with probability $1 - P$
- This means that our model has one parameter - P
- Suppose we observed 100 tosses and 53 of them gave heads
- What is P ?

Coin tossing example

Frequentist answer (aka maximum likelihood)

Pick the value of p (probability of head) that makes the observation A of 53 heads and 47 tails most probable.

$$P(A) = p^{53} (1 - p)^{47}$$

Differentiate over p

$$\frac{dP}{dp} = 53p^{52}(1 - p)^{47} - 47p^{53}(1 - p)^{46} = \left(\frac{53}{p} - \frac{47}{1 - p} \right) [p^{53}(1 - p)^{47}] = 0$$

$$\frac{53}{p} = \frac{47}{1 - p}$$

$$53 - 53p = 47p$$

$$p = 0.53$$

Coin tossing example

Bayesian answer

- Start with a prior distribution over p :
$$P(p) = 1$$
- Step 1a: we get a head after the first toss, multiply by the probability of observing a head given that value
- Step 1b: rescale probability density to 1
- Step 2a: we get a tail after the second toss, multiply by the probability of observing a tail given that value
- Step 2b: rescale probability density to 1

Coin tossing example

Bayesian answer

- Start with a *prior* distribution over p :
$$P(p) = 1$$
- Step 1a: we get a head after the first toss, multiply by the probability of observing a head given that value
- Step 1b: rescale probability density to 1
- Step 2a: we get a tail after the second toss, multiply by the probability of observing a tail given that value
- Step 2b: rescale probability density to 1
- ...
- After 100 steps we get a *posterior* probability distribution

ML models

- Probabilistic modelling such as Naïve Bayes algorithm which ML classifier based on applying Bayes' theorem while assuming that the features in the input are all independent
- Kernel methods such support vector machines (SVM)
- Decision trees, random forests an gradient boosting machines
- K-nearest neighbours

Parametric vs Non-parametric models

Parametric model has a fixed number of parameters

E.g., linear regression

Non-parametric model has a number of parameters which grows with the amount of training data.

E.g. K-nearest neighbour (KNN)

Deep Learning

- DL is strongly related to mathematical statistics
- but the statistical apparatus has to be much simplified when faced with computational complexity and the amount of data (e.g. millions of images)
- Therefore DL exhibits relatively little mathematical theory and is engineering oriented
- Often ideas are proposed based on (sometimes limited) empirical evidence
- Nevertheless there has been incredibly good progress in the last couple of years



Learning process

Here is Madonna.
Assume you didn't know who she is



This is your training set.

Now, do you recognise these people?

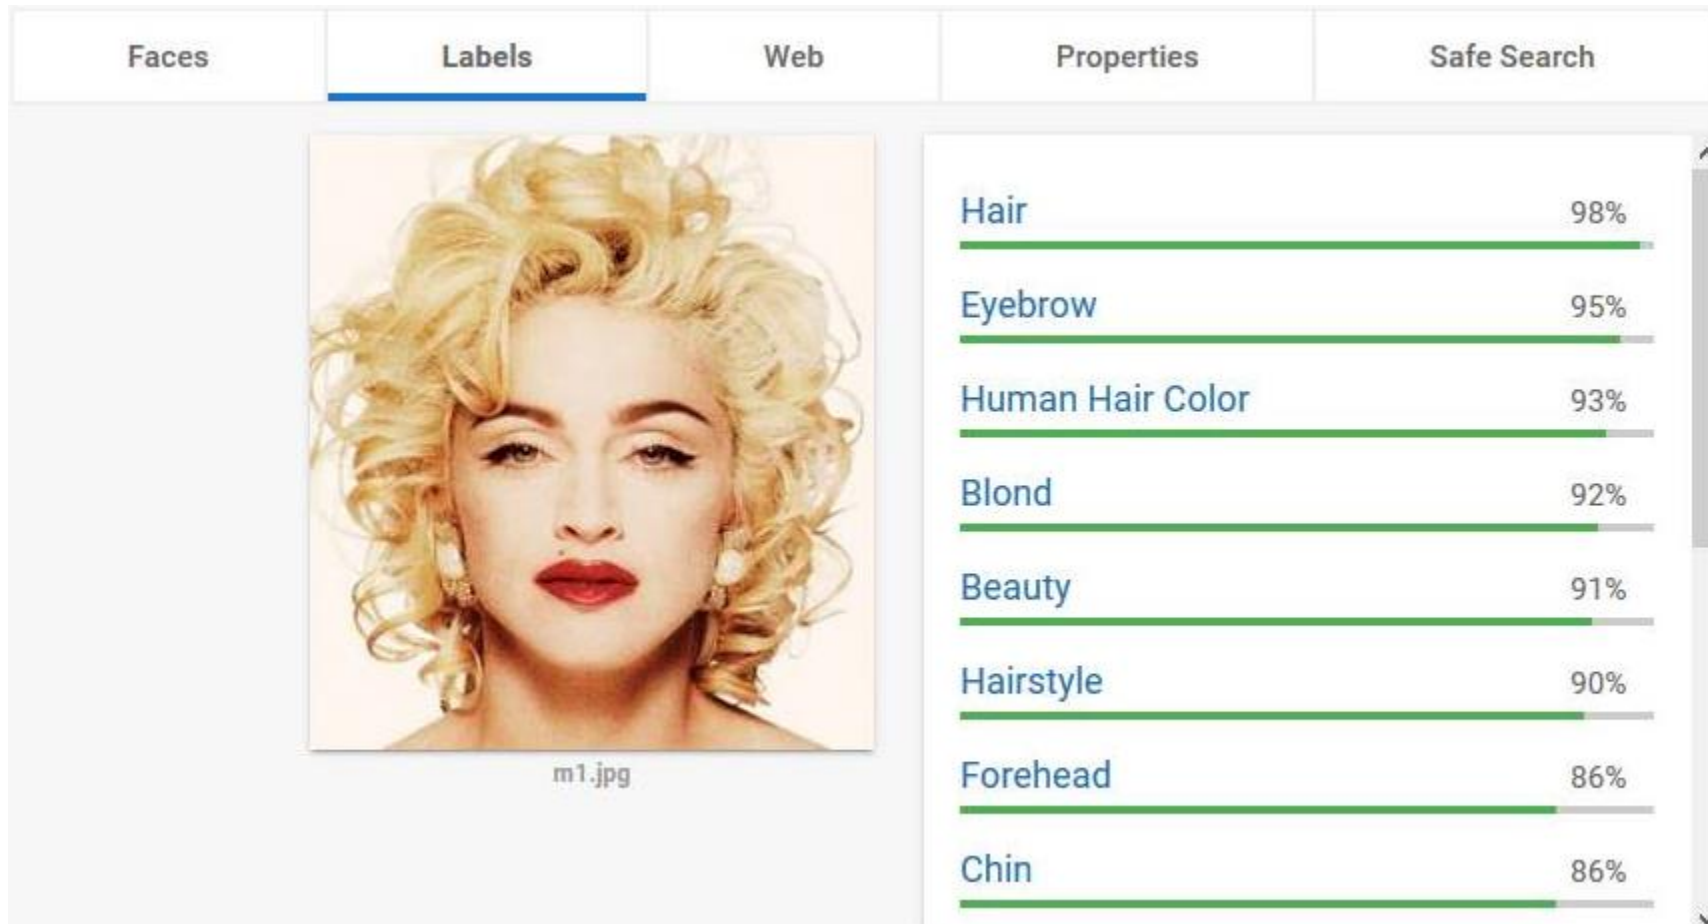
45



This is your test set.

What about ML? Can a machine do it?

<https://cloud.google.com/vision/>




Not quite what we wanted to see

Can a machine recognise image?

<https://cloud.google.com/vision/>

FacesLabelsWebPropertiesSafe Search



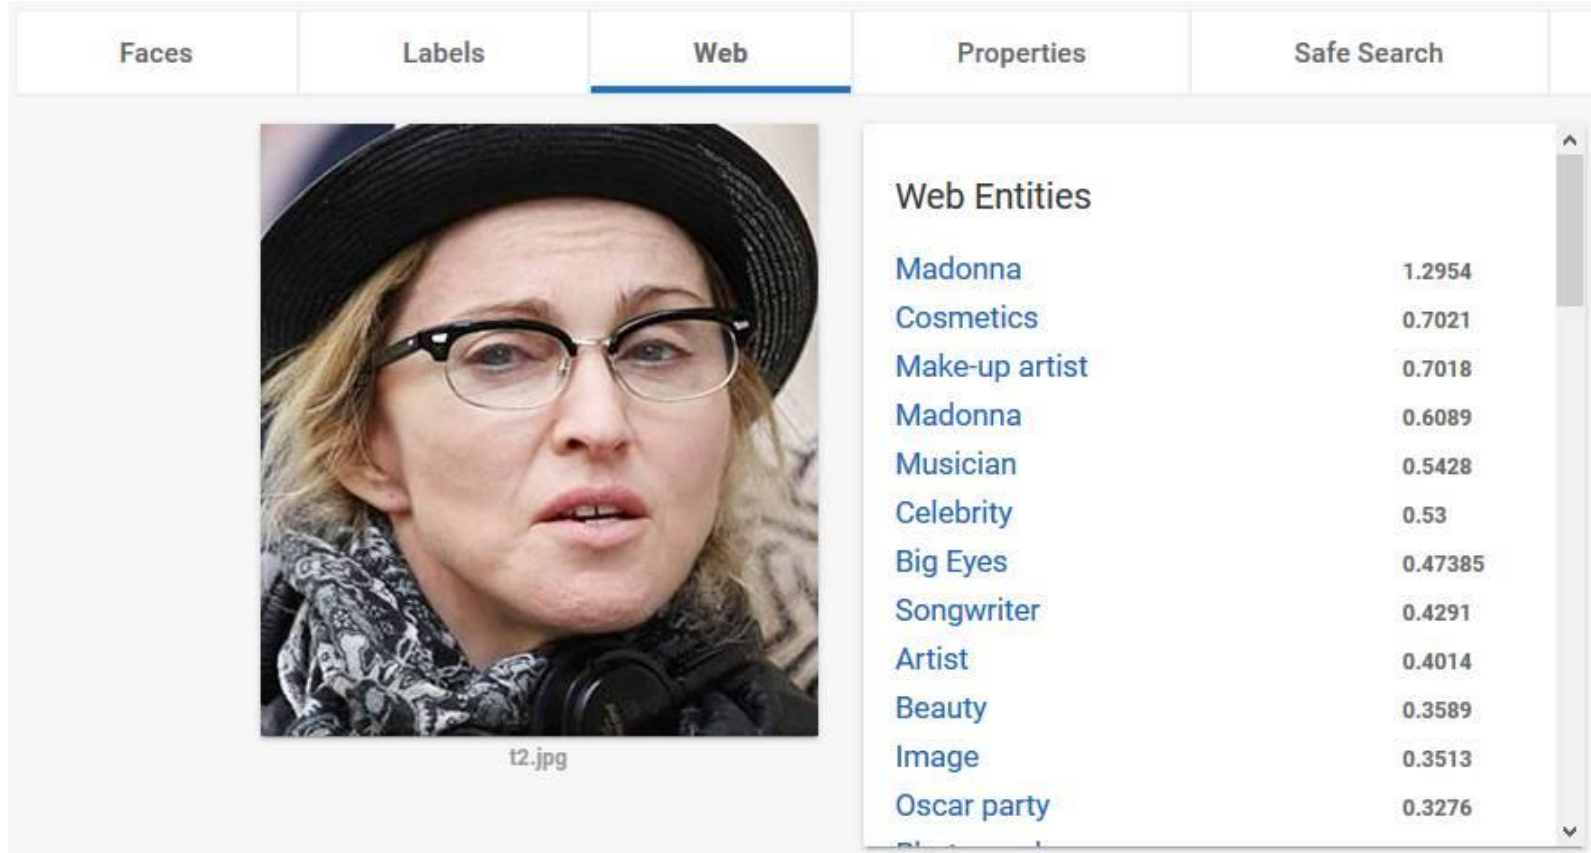
m1.jpg

Web Entities

Madonna	12.8475
Blond Ambition World Tour	0.78135
Singer	0.672
	0.6447
1990s	0.63825
Madonna	0.5916
Music	0.5844
	0.5647
Photography	0.5331
Female	0.5178
Vogue	0.5
Vogue	0.4494
Cleavage	0.4101

Can a machine recognise image?

<https://cloud.google.com/vision/>



The screenshot displays the Google Cloud Vision API interface. At the top, there are five tabs: 'Faces', 'Labels', 'Web', 'Properties', and 'Safe Search'. The 'Web' tab is currently selected. Below the tabs, on the left, is a photograph of Madonna wearing a black hat and glasses. Below the photo is the filename 't2.jpg'. To the right of the photo is a list of 'Web Entities' with their corresponding confidence scores.

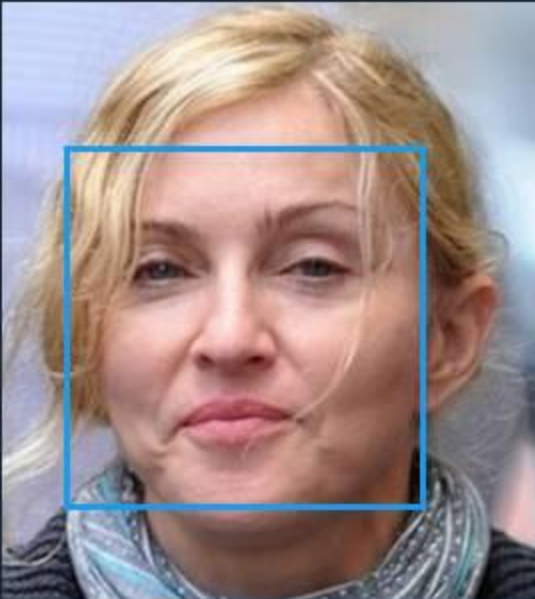
Web Entity	Score
Madonna	1.2954
Cosmetics	0.7021
Make-up artist	0.7018
Madonna	0.6089
Musician	0.5428
Celebrity	0.53
Big Eyes	0.47385
Songwriter	0.4291
Artist	0.4014
Beauty	0.3589
Image	0.3513
Oscar party	0.3276

If we use the first web entity as a predictor then Cloud Vision API identified all the test images correctly. However this should not come as surprise because the images were collected through Google search in the first place.

Can a machine recognise image?

<https://clarifai.com/demo>

MORE MODELS ▼



Celebrity

PREDICTED CONCEPT	PROBABILITY
Madonna	0.405


Clarifai was only able to guess the first test image with $P = 0.4$. It did not recognise the second “training” image ($P(m1) = 0.989$ and $P(m3) = 0.543$).

Can a machine recognise image?

Amazon Rekognition

Celebrity recognition

Rekognition automatically recognizes celebrities in images and provides confidence scores.



Done with the demo?

[Learn more](#)

▼ Results



Shep Pettibone

[Learn More](#)

Match confidence

100 %

► Request

► Response

Amazon has special Celebrity recognition mode. Unfortunately it screwed up badly ($P(m1) = 1$, $P(m3) = 0.96$ and nothing else at all).


Can a machine compare images?

Amazon Rekognition


Face comparison

Compare faces to see how closely they match based on a similarity percentage.

Reference face



Comparison faces



Done with the demo?
[Learn more](#)

▼ Results

 $=$ 

Similarity 94 %

► Request


► Response

Rekognition also has face comparison mode and at least it is self-consistent. But that's pretty much it – none other pair of images is comparable.

Just to show that Amazon Rekognition may work quite well

Text in image

Rekognition automatically detects and extracts text in your images. [Learn More](#)



Done with the demo?

[Learn more](#)

▼ Results

US English only

EAPEN | 05/07/2013 | PHILIPS |
 101153 | FR.MULLER | HOSPITAL |
 MANGALORE | 10:15:50 |
 LIVER |
 -ABDOMEN | N |
 C5-2 |
 MI | 1.4 |
 TIS | 0.9 |
 RT | ATRIAL | LESION | -H3 | Gn |
 100 |

Amazon Rekognition in text detection mode.

Note, that examples above were Demos out-of-the-box, without adjustments or tuning.

Applications of ML by industry



HEALTHCARE

Making medical diagnoses more quickly and more accurately.



SOCIAL MEDIA

Customizing content based on user behavior, such as in your Facebook feed.



RETAIL

Recommending products based on past behavior and similar customers.



MANUFACTURING

Anticipating repairs and improving preventive maintenance.



MARKETING

Analyzing customer responses to ads.



SECURITY

Fending off cyber attacks based on anomalous behavior.



TRANSPORTATION

Analyzing & responding to the real-world environment in the development of driverless cars.



CUSTOMER SERVICE

Creating intelligent virtual bots to manage customer interactions.



REAL ESTATE

Generating property recommendations.



GAMING

Analyzing competitive play to anticipate moves and create more challenging enemies.

ML may be not perfect but it is coming

- The logic behind the new revolution is that this time it's prediction which is going to become cheap
- Before the breakthroughs were computing, then internet search and now prediction
- The challenge is to build (reliable) systems
- The foundation has been laid out (or so we are lead to believe)
- Just as it is very useful to know how search or computing work, it is good to know how predictions work even if you are not planning to create ML models yourself
- Think of the OCF/IBM ML demo which analysed a medical dataset without asking or telling what the model was
- We are likely to see more tools with ML engines inside in the coming years

Why Deep Learning works?

- The world is **hierarchical**
- **Manifold** hypothesis - most input configurations are unlikely, real examples concentrate near a (much) lower dimensional manifold in feature space
- Although a single layer can be as powerful as **many layers**, deep (nested) layers need fewer neurons to create powerful representations
- Layer action is mostly linear (kernels are simple computationally)
- Which helps finding derivatives used in **backpropagation**
- Increased computational power (clusters of GPUs, TPUs)
- Availability of training data
- Technological advances such as batch normalisation, dropout etc
- Stochastic Gradient Descent (SGD); global minimum is not essential; optimisation process is never fully converged
- Human perception is fast: ~ 0.1 sec.
Neurons fire at most 100 times a sec.
Hence our neurons fire 10 times at most.
Therefore 10 layers can solve any perception task

Tom Goldstein –

What do neural loss surfaces look like?



Typical Machine Learning Workflows

Traditional Machine Learning: Given data, **design feature set** and use a suitable ML algorithm to train a model, test on test data set.

Typical Deep Learning: Given data, **design DL architecture** and use it to train a model. Feature sets are created automatically (trained).


1. train (iteratively if necessary) on training set,
2. validate on validation set (change hyper parameters if necessary and repeat),
3. test on test set (this is a test of the power of your model: its accuracy and generalisation)

Full Dataset:


Training Data	Validation Data	Test Data
---------------	-----------------	-----------

Deep Learning uses Neural Networks

which are inspired by nature



Simple Three Layer Neural Network




This is a flat network. Deep networks have more layers or network hops.

Output layer - layer 3


Hidden layer - layer 2

Input layer - layer 1



Hierarchical multi-layer network

60



Training on 10 M
unlabelled frames from
YouTube videos
(unsupervised learning).

Hierarchy of Representations

Instead of traditional approach using feature design, thanks to emergence of large data sets, features are trained directly from data (which is why number of parameters in models is quite large – tens of millions is common). Input is pre-processed (e.g. average subtracted, signal normalised)

From concrete to abstract:

- vision: pixel, motif, part, object
- text: character, word, clause, sentence
- speech: audio, band, phoneme, word

MNIST - one of the first successes


‘Hello World’ of Deep Learning

- created by Yann LeCun et al
- database of handwritten digits
- a training set of 60,000 and a test set of 10,000 labelled 28x28 images
- end-of-life as a useful benchmark

Type	Error Rate %
Linear Classifier	7.6
K-Nearest Neighbours	0.52
Support vector machine	0.56
Convolutional neural network	0.23




Progress in recent years - Visual Object Classification



Convolutional Networks

64

- Features are discovered automatically; layer 1 features resemble Gabor filters; layers represent hierarchy of features



- Training is done iteratively (forward pass is followed by backprop) and may take a long time
- Once the training is done the inference (the forward pass) is usually quick making real time tagging (identification of live stream) possible.

ConvNet

Find probabilities (or scores) of labels



Maximally accurate	Maximally specific
cat	1.80727
domestic cat	1.74727
feline	1.72787
tabby	0.99133
domestic animal	0.78542

Semantic Image Segmentation

66



Objects appearing in the image:
Bicycle,
Person



Image captions

67

Describes without errors



A person riding a motorcycle on a dirt road.



A group of young people playing a game of frisbee.



A herd of elephants walking across a dry grass field.

Describes with minor errors



Two dogs play in the grass.



Two hockey players are fighting over the puck.



A close up of a cat laying on a couch.

Somewhat related to the image



A skateboarder does a trick on a ramp.



A little girl in a pink hat is blowing bubbles.



A red motorcycle parked on the side of the road.

Unrelated to the image



A dog is jumping to catch a frisbee.

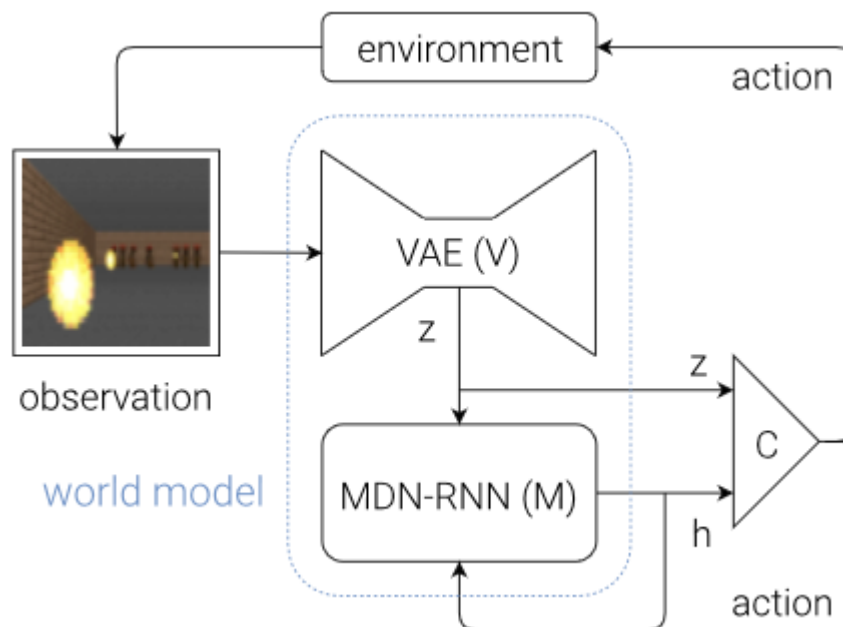


A refrigerator filled with lots of food and drinks.



A yellow school bus parked in a parking lot.

Recent headlines



World model to play Doom (reinforcement learning)



Progressive Growing of GANs

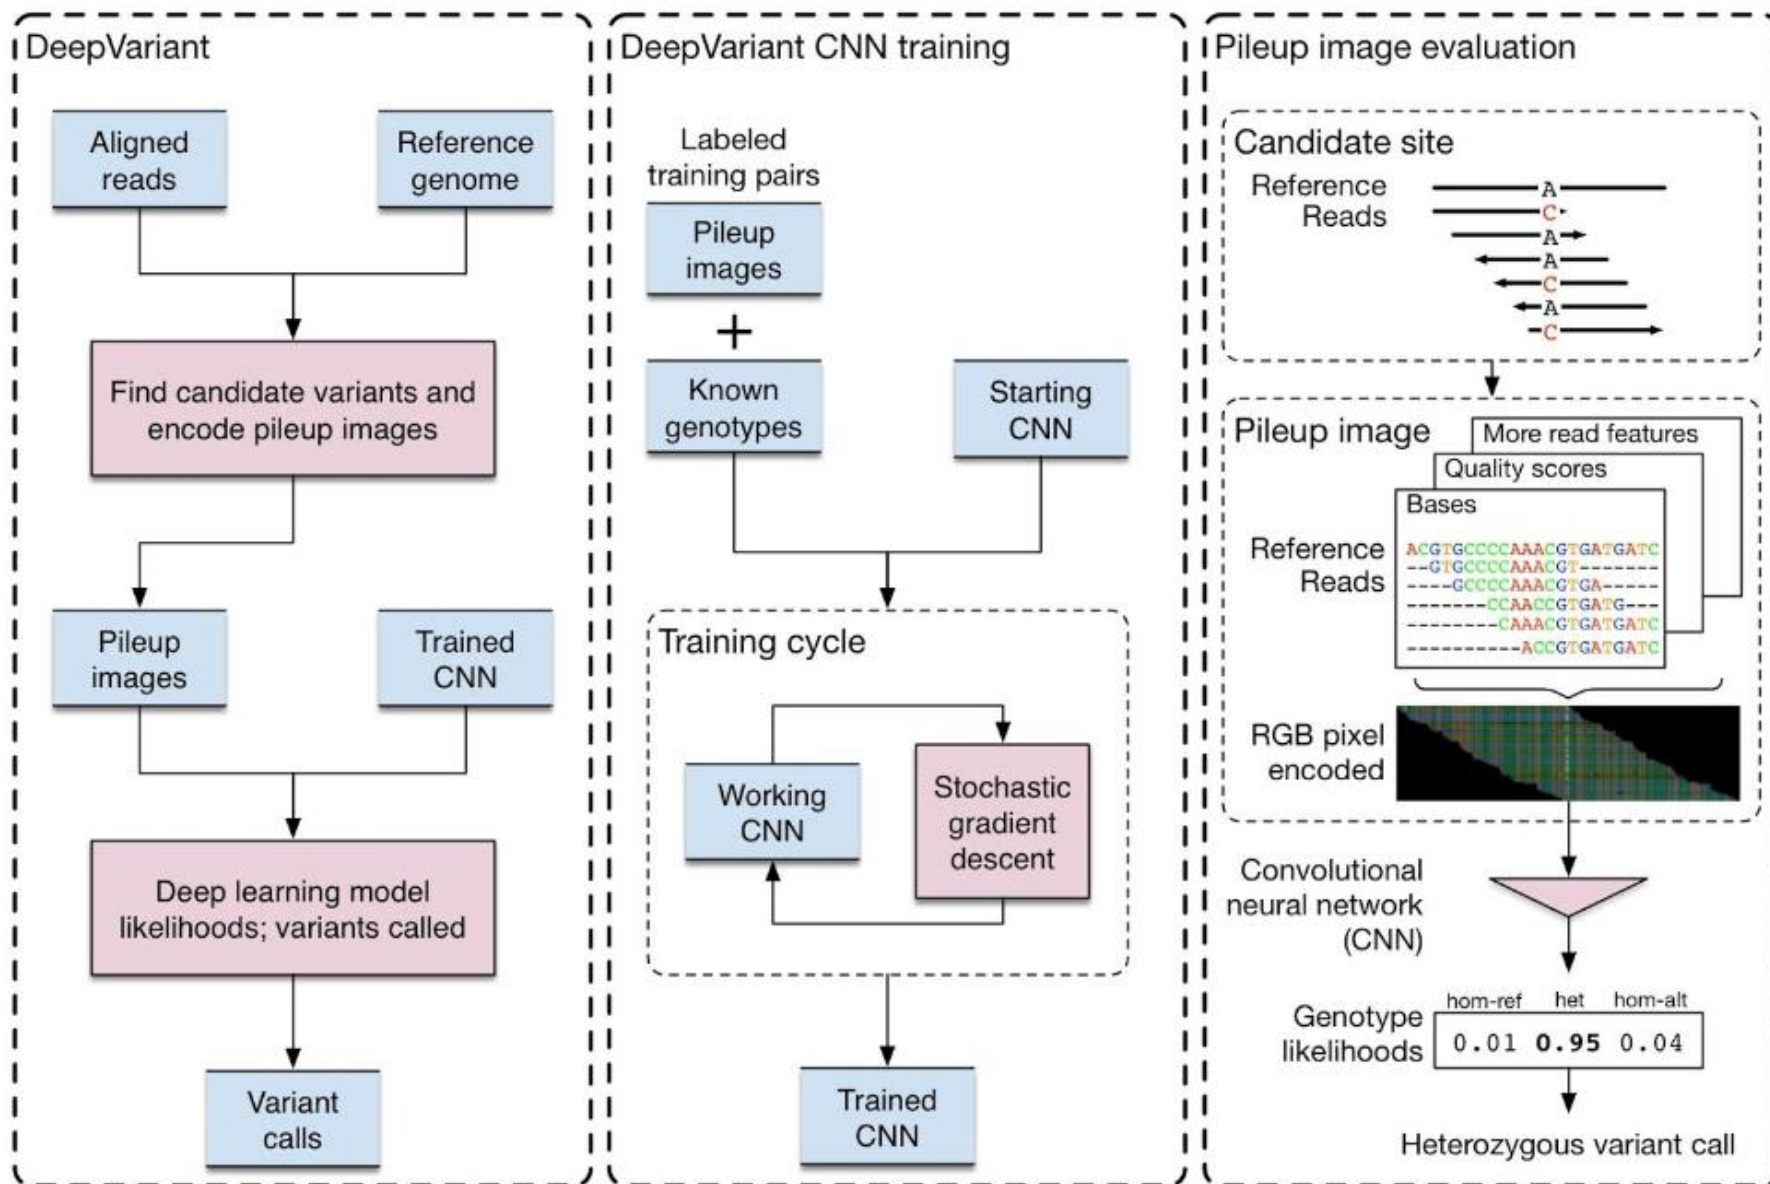
https://github.com/tkarras/progressive_growing_of_gans



From AlphaGo to AlphaZero (DeepMind)

DeepVariant

69



References and homework

- François Chollet, *Deep Learning with Python*, Manning 2017
Chapter 1: What is Deep Learning
<https://www.manning.com/books/deep-learning-with-python>
- Ian Goodfellow and Yoshua Bengio and Aaron Courville, *Deep Learning*, MIT Press 2016
Chapter 5: Machine Learning Basics
<http://www.deeplearningbook.org/contents/ml.html>
- Homework to refresh Numpy
https://notebooks.azure.com/ink1/libraries/IntroductionToPython/html/numpy_intro.ipynb
- Let us know if you would like to join Deep Learning group at ICR
<https://nexus.icr.ac.uk/teams/Deep%20Learning/>

Next week Sebastian Poelsterl will cover

1. Linear Algebra

- Scalar, Vector, Matrix and Tensor
- Transpose
- Dot product
- Identity matrix
- Linear system of equations
- Matrix inverse
- Matrix and vector norms
- Symmetric and orthogonal matrix

2. Calculus

- Gradient
- Hessian / curvature
- Gradient descent / iterative optimisation
- Minimum / Maximum / Saddle Point

3. Least squares

4. Logistic regression

For refresher, see <https://www.khanacademy.org/math/linear-algebra>

ICR

